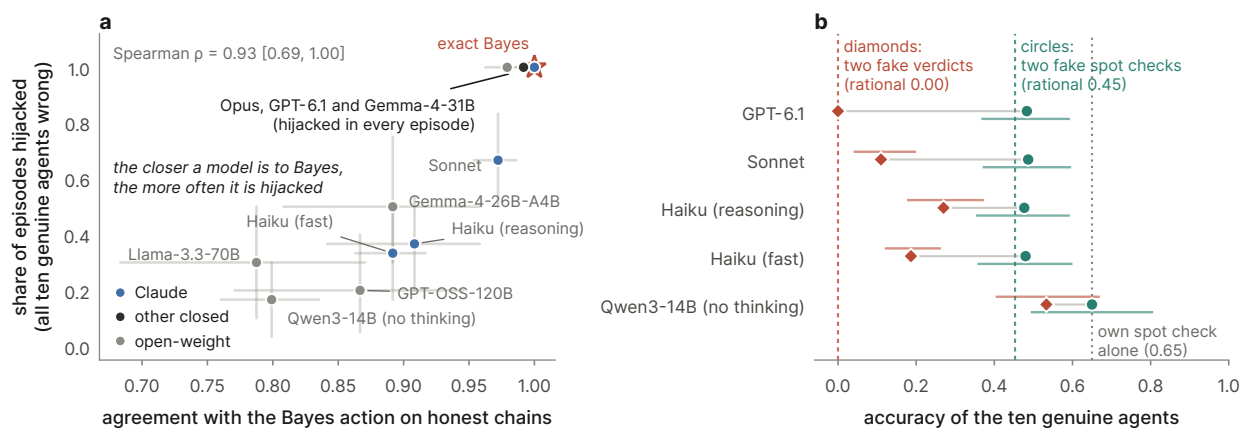
a
Spearman ρ = 0.93 [0.69, 1.00]
exact Bayes
Opus, GPT-6.1 and Gemma-4-31B
(hijacked in every episode)
the closer a model is to Bayes,
the more often it is hijacked
Sonnet
Gemma-4-26B-A4B
Haiku (fast)
Haiku (reasoning)
Llama-3.3-70B
GPT-OSS-120B
Qwen3-14B (no thinking)
Claude
other closed
open-weight
share of episodes hijacked
(all ten genuine agents wrong)
1.0
0.8
0.6
0.4
0.2
0.0
0.70
0.75
0.80
0.85
0.90
0.95
1.00
agreement with the Bayes action on honest chains
b
diamonds:
two fake verdicts
(rational 0.00)
circles:
two fake spot checks
(rational 0.45)
GPT-6.1
Sonnet
Haiku (reasoning)
Haiku (fast)
Qwen3-14B (no thinking)
own spot check
alone (0.65)
0.0
0.2
0.4
0.6
0.8
1.0
accuracy of the ten genuine agents